

# Course Overview

LECTURE 5 - part I

Knowledge Representation & Logics  
Logical Arguments • Propositional Logic  
Validity, Equivalence, Satisfiability, Entailment  
Inference by Natural Deduction



**UNSW**  
SYDNEY

# Planned Topics

## AI, Agents & AI Programming

- What is AI?
- Agents, Agent Types
- Agents Tasks
- AI Programming

## Machine Learning

- Supervised Machine Learning
- Perceptrons & Neural Networks
- Learning to Act - Reinforcement Learning

## Solving Problems by Search

- Path Search – Basic search Algorithms
- Informed - Heuristic Path Search
- Constraint Satisfaction Problems

## Knowledge and Reasoning

- Planning
- Propositions and Inference
- Reasoning with Uncertainty
- Knowledge Representation

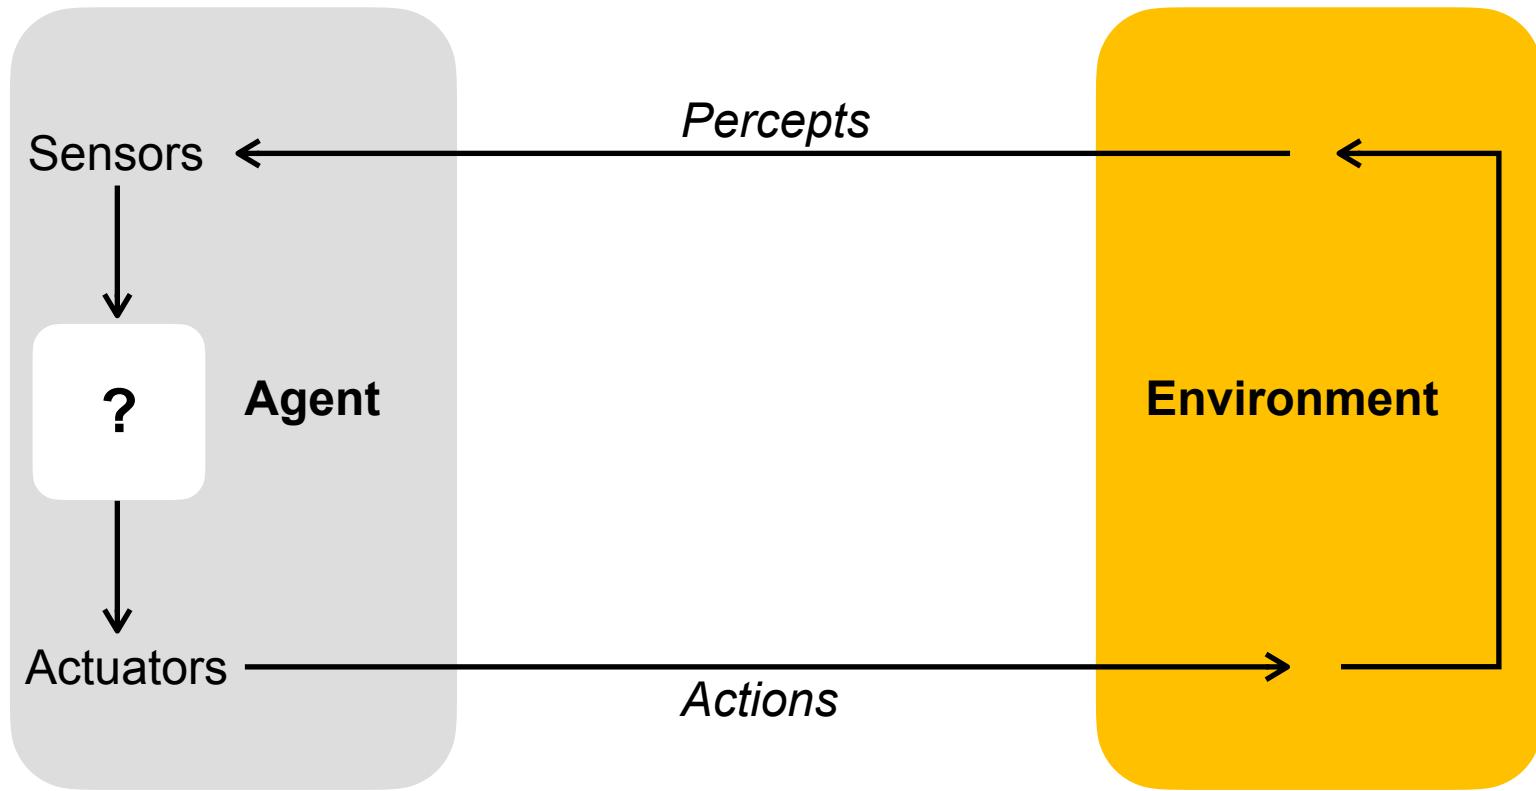
# Assessment

- Participation: 10%
- Assignments: 40%
  - assignment 1: 20%
  - assignment 2: 20%
- Exam: 50% + extra points

**In order to pass the course, you must score:**

- at least 20/40 for the assignments
- at least 20/50 for the exam
- a combined mark of at least 50/100

# Agent Model



# Classifying environment types

Simulated vs Situated/Embodied

Static vs Dynamic

Discrete vs Continuous

Fully Observable vs Partially Observable

Deterministic vs Stochastic

Episodic vs Sequential

Known vs Unknown

Single-Agent vs Multi-Agent

# Environment Types

**Simulated:** a separate program is used to simulate an environment, feed percepts to agents, evaluate performance, etc.

**Static:** environment doesn't change while the agent is deliberating

**Discrete:** finite (or countable) number of possible percepts/actions

**Fully Observable:** percept contains all relevant information about the world

**Deterministic:** current state of world uniquely determines the next

**Episodic:** every action by the agent is evaluated independently

**Known:** the rules of the game, or physics/dynamics of the environment are known to the agent

**Single-Agent:** only one agent acting in the environment

# Specifying Agents

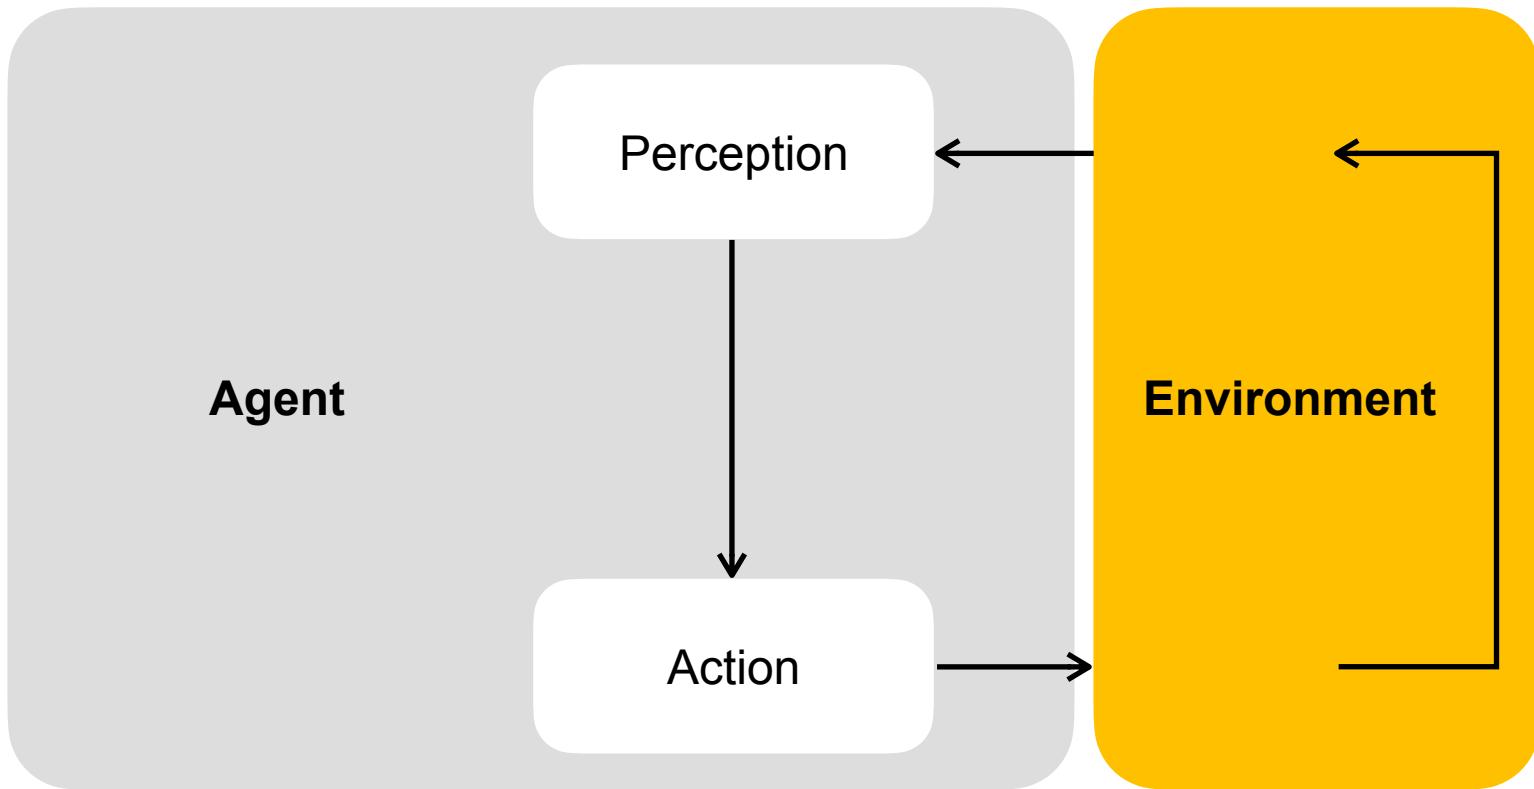
- **percepts**: inputs to the agent via sensors
- **actions**: outputs available to the agent via effectors
- **goals**: objectives or performance measure of the agent
- **environment**: world in which the agent operates

Most generally, a function from percept sequences to actions

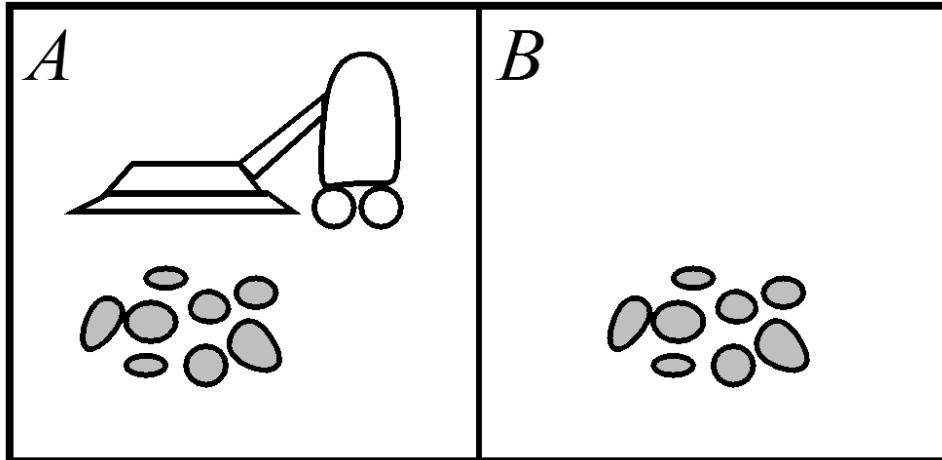
Ideally **rational agent** does whatever action is expected to maximise some performance measure – the agent may not know the performance measure (Russell and Norvig 2010)

**Resource bounded** agent must make “good enough” decisions based on its perceptual, computational and memory limitations (design tradeoffs)

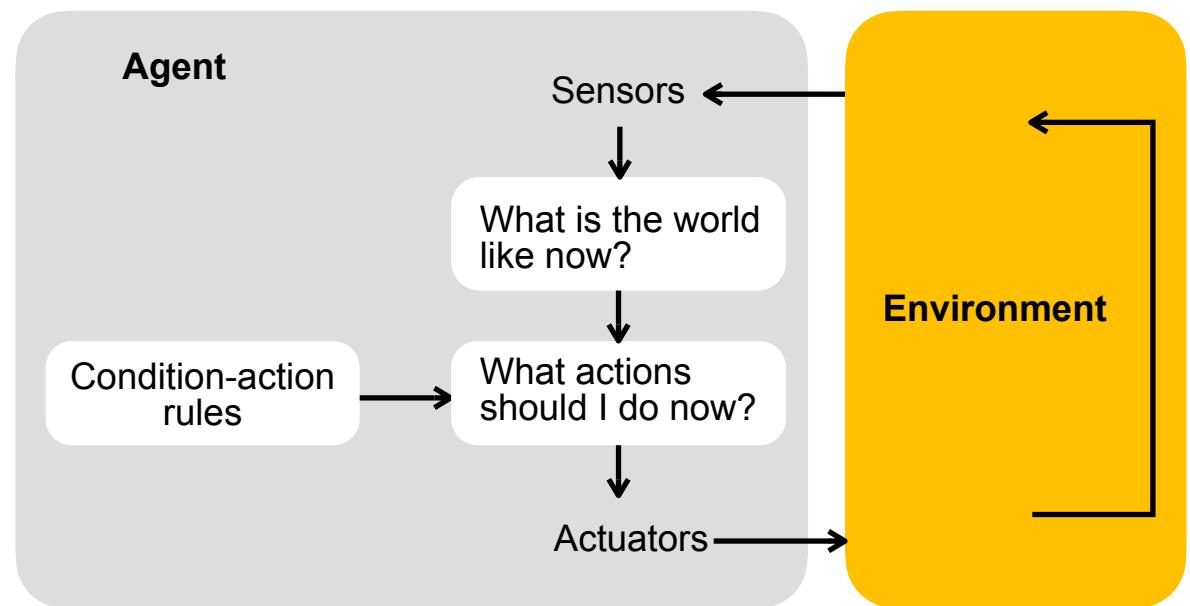
# Reactive agent



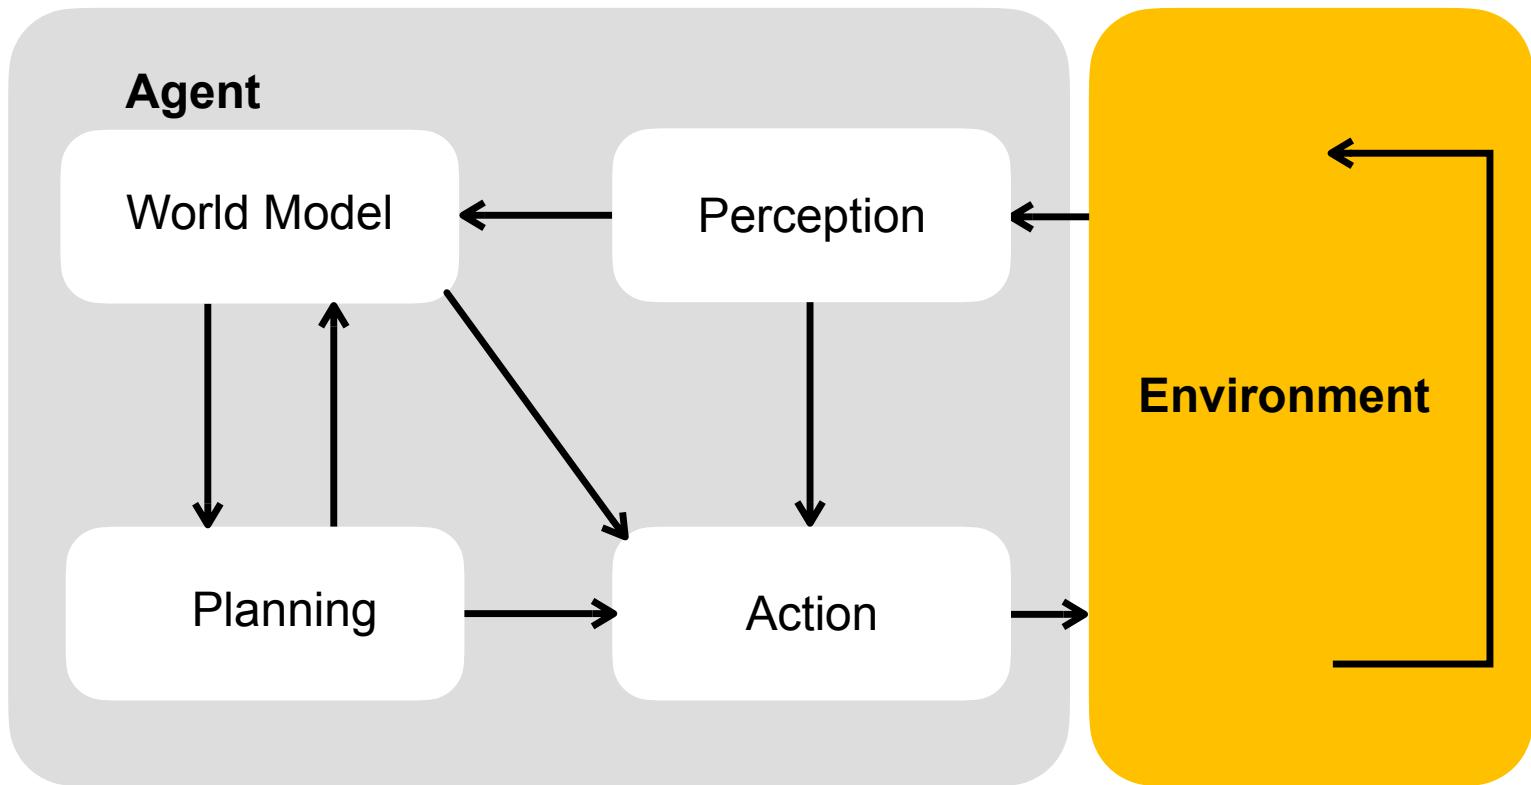
# Vacuum-cleaner world



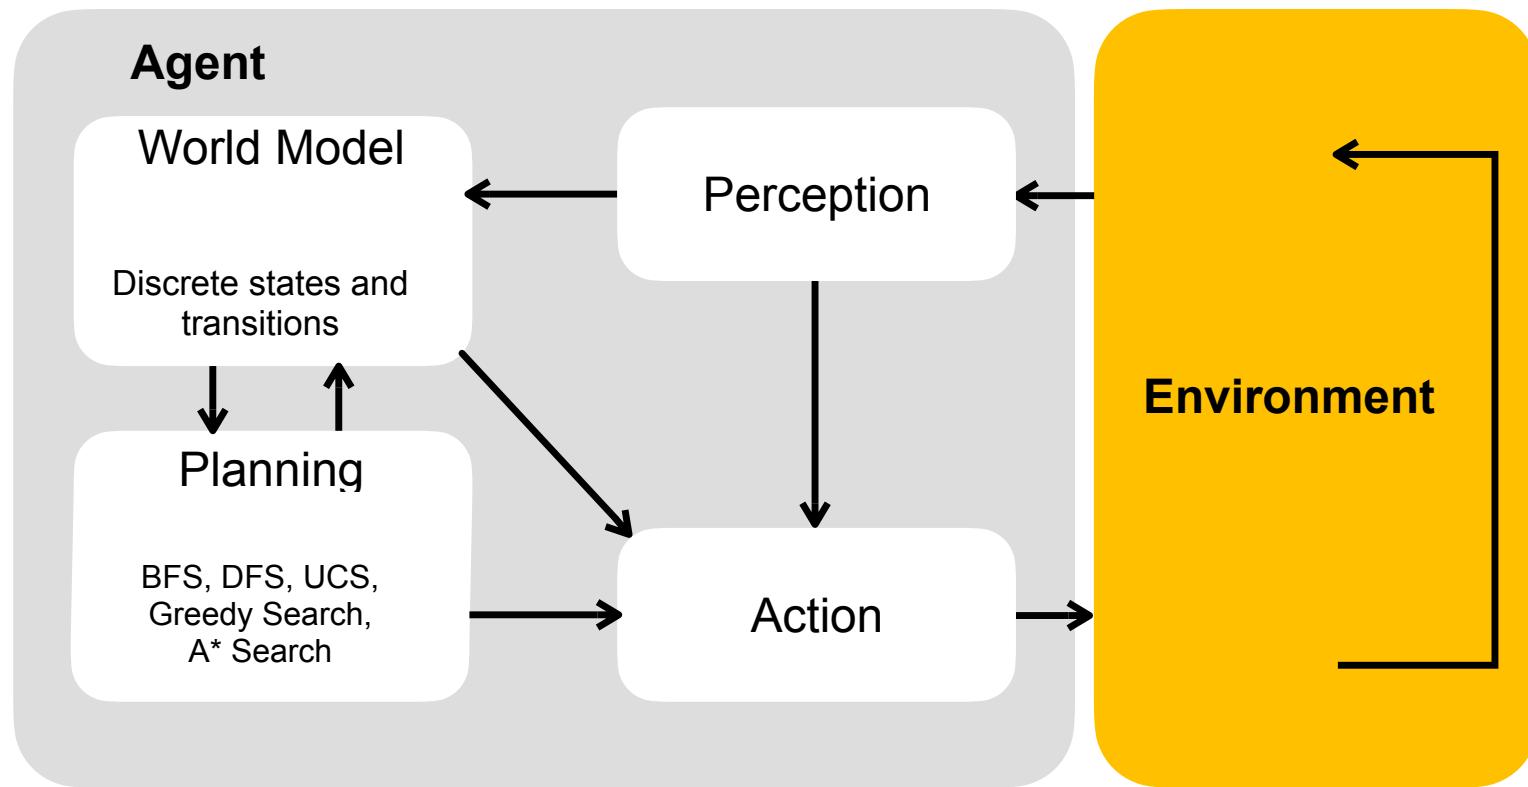
Percepts: location and contents  
Actions: Left, Right, Suck, NoOp



# Planning Agent



# Path Search Agent



# State space problem

A state-space problem consists of:

- a set of states
- a subset of states called the **start states**
- a set of actions
- an **action function**: given a state and an action, returns a new state
- a (set of) goal state, specified as Boolean function, goal(s)
- a criterion that specifies the quality of an acceptable solution.

# Single-state Space Specification

A task is specified by states and actions:

- state space e.g. other rooms
- initial state e.g., “outside of room 103”
- actions or operators (or successor function  $S(x)$ ) = set of action–state pairs
- goal test, check if a state is goal state
  - In this case, there is only one goal specified (“inside of room 123”)
- path cost e.g. sum of distances, number of actions etc.
- a solution is a sequence of actions leading from the initial state to a goal state

# Assignment 2 - State Space

```
class GameFifteenProblem(Search_problem):

    def __init__(self, start, goal):
        return

    def start_node(self):
        """Returns the start node"""
        return

    def is_goal(self, node):
        """Returns True if the node is the goal, otherwise False"""
        return

    def neighbors(self, node):
        """Returns a list of the arcs for the neighbors of node, for example:
        return [Arc(node, to_neighbor_node1, cost1), Arc(node, to_neighbor_node2, cost2)]"""
        return

    def heuristic(self, node):
        """Returns the heuristic value of the node
        based on the Manhattan distance"""
        return
```

# Assignment 2 - Heuristics

```
class GameFifteenProblemEuclidean(GameFifteenProblem):
    def __init__(self, start, goal):
        super().__init__(start, goal)

    def heuristic(self, node):
        """Returns the heuristic value of the node
        based on the Euclidean distance"""
        return
```

```
class GameFifteenProblemInversions(GameFifteenProblem):
    def __init__(self, start, goal):
        super().__init__(start, goal)

    def heuristic(self, node):
        """Returns the heuristic value of the node
        based on the sum of the inversion number of a permutation"""
        return
```

# Assignment 2 - State Space

Heuristic	h-value: start14	h-value: start17	h-value: start23
-----------	------------------	------------------	------------------

Manhattan	#	#	#
Euclidean	#	#	#
Inversions	#	#	#

Heuristic	expanded: start14	expanded: start17	expanded: start23
-----------	-------------------	-------------------	-------------------

Manhattan	#	#	#
Euclidean	#	#	#
Inversions	#	#	#

Heuristic	path cost: start14	path cost: start17	path cost: start23
-----------	--------------------	--------------------	--------------------

Manhattan	#	#	#
Euclidean	#	#	#
Inversions	#	#	#

# Path Search Algorithms - Search Strategies

General Search algorithm:

- add initial state to queue
- repeat:
  - take node from front of queue
  - test if it is a goal state; if so, terminate
  - “expand” it, i.e. generate successor nodes and add them to the queue

Search strategies are distinguished by the order in which new nodes are added to the queue of nodes awaiting expansion.

# Search Strategies

- **BFS** and **DFS** treat all new nodes the same way:
  - BFS add all new nodes to the back of the queue (**FIFO**)
  - DFS add all new nodes to the front of the queue (**FILO**)
- **Best First Search** uses an evaluation function  $f()$  to order the nodes in the queue;
  - Similar to **UCS**  $f(n) = g(n)$  of path from root to node  $n$
- Informed or Heuristic search strategies incorporate into  $f()$  an estimate of distance to goal
  - **Greedy Search**  $f(n) = \text{estimate } h(n)$  of cost from node  $n$  to goal
  - **A\* Search**  $f(n) = g(n) + h(n)$

# Search Strategies

- A strategy is defined by picking **the order of node expansion**
- Strategies are evaluated along the following dimensions:
  - completeness – does it always find a solution if one exists?
  - time complexity – number of nodes generated/expanded
  - space complexity – maximum number of nodes in memory
  - optimality – does it always find a least-cost solution?
- Time and space complexity are measured in terms of
  - $b$  – maximum branching factor of the search tree
  - $d$  – depth of the least-cost solution
  - $m$  – maximum depth of the state space (may be  $\infty$ )

# Complexity Results for Search Algorithms

Algorithm	Completeness	Admissibility	Space	Time
Breadth-First Search	guaranteed, if $b$ is finite	guaranteed, if arcs have the same cost	$O(b^d)$	$O(b^d)$
Depth-First Search	not guaranteed	not guaranteed	$O(bm)$	$O(b^m)$
Uniform Cost Search	guaranteed, if $b$ is finite and $\text{cost} \geq \varepsilon$	guaranteed	$O(b^{[1 + C/\varepsilon]})$	$O(b^{[1 + C/\varepsilon]})$
Depth Limited Search	guaranteed, if $b$ is finite	not guaranteed	$O(bk)$	$O(bk)$
Iterative Deepening Search	guaranteed, if $b$ is finite	guaranteed, if arcs have the same cost	$O(bd)$	$O(b^d)$
Bidirectional Search	depends on search algorithms used	depends on search algorithms used	$O(b^{d/2})$	$O(b^{d/2})$
Greedy Best-First Search	guaranteed, if $b$ is finite and with repeated-state checking	not guaranteed	$O(b^m)$	$O(b^m)$
A* Search	guaranteed	guaranteed if heuristic is admissible/ consistent	$O(b^d)$	$O(b^d)$
Iterative Deepening A* Search	guaranteed	guaranteed if heuristic is admissible/ consistent	$O(d)$	$O(b^d)$

# Assignment 2 - complexity vs reality

## Part 3: Deceptive Starting States [3 marks]

### Task 3.1

Run IDA\* on the starting states below and report the number of the expanded nodes.

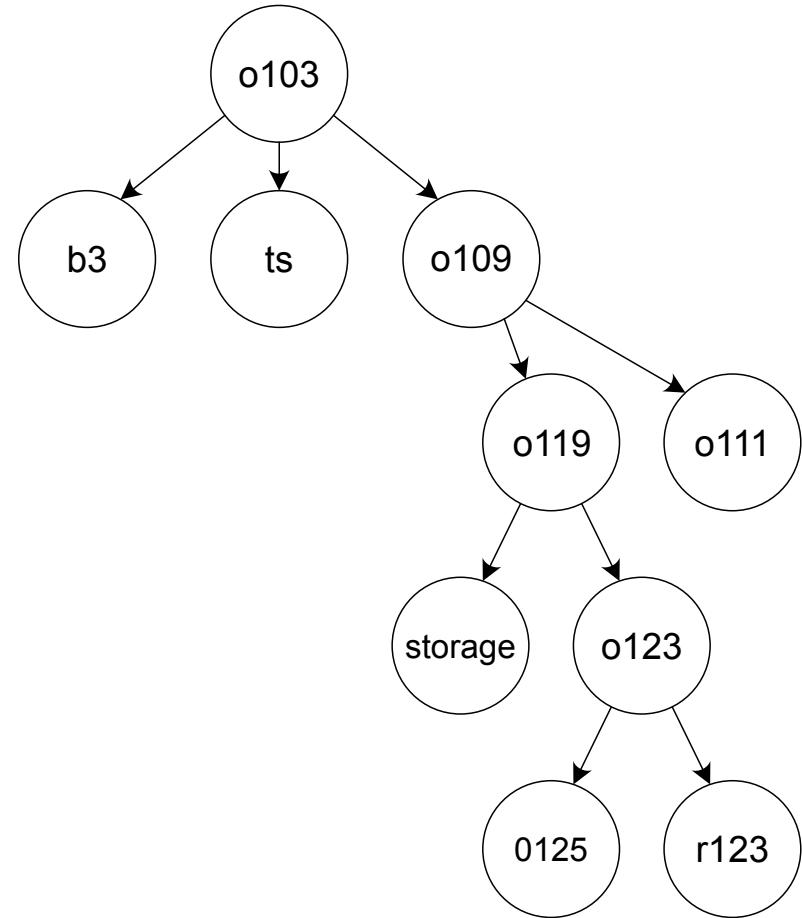
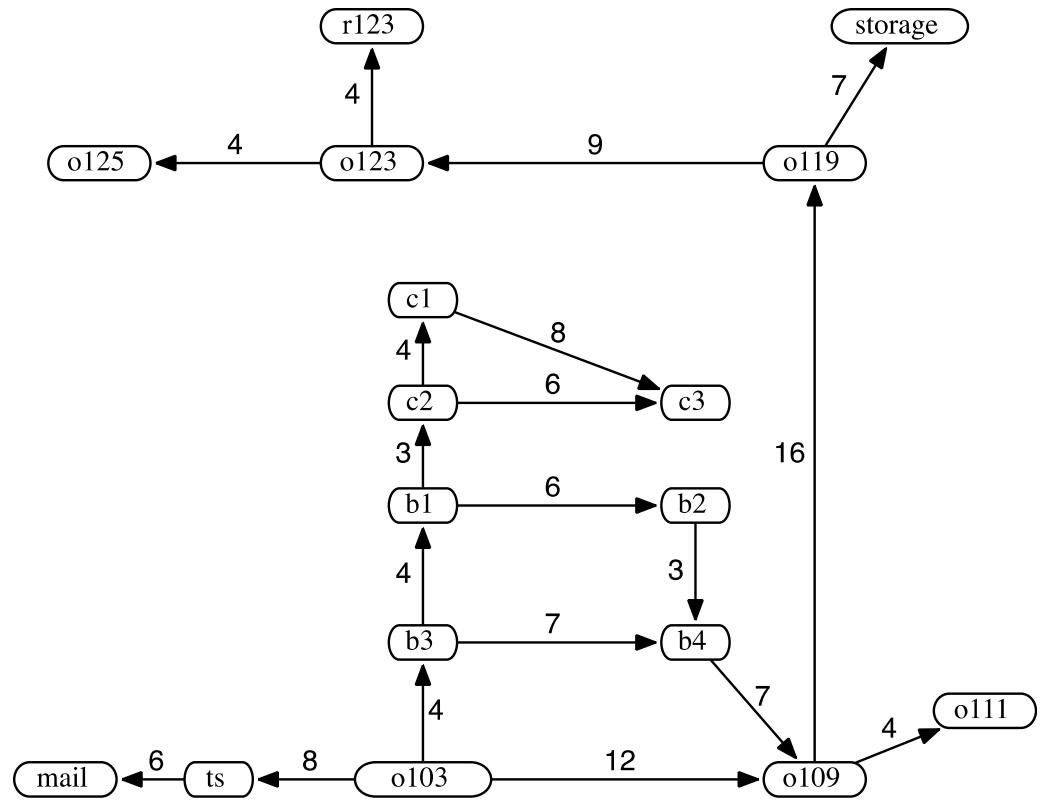
```
start37_1 = [[7, 11, 12, 4],  
             [2, 3, 14, 1],  
             [6, 5, 13, 0],  
             [9, 10, 15, 8]]  
  
start37_2 = [[7, 11, 12, 4],  
             [2, 3, 8, 14],  
             [6, 0, 1, 15],  
             [9, 5, 10, 13]]  
  
# your code
```

### Task 3.2

Explain why there is a difference between the number of the expanded nodes for these starting states if their costs of the optimal paths are the same.

```
# your answer
```

# DFS



# Constraint Satisfaction Problems (CSPs)

- Constraint Satisfaction Problems are defined by a set of variables  $X_i$ , each with a domain  $D_i$  of possible values, and a set of constraints  $C$ .
- The aim is to find an assignment of the variables  $X_i$  from the domains  $D_i$  in such a way that none of the constraints  $C$  are violated.

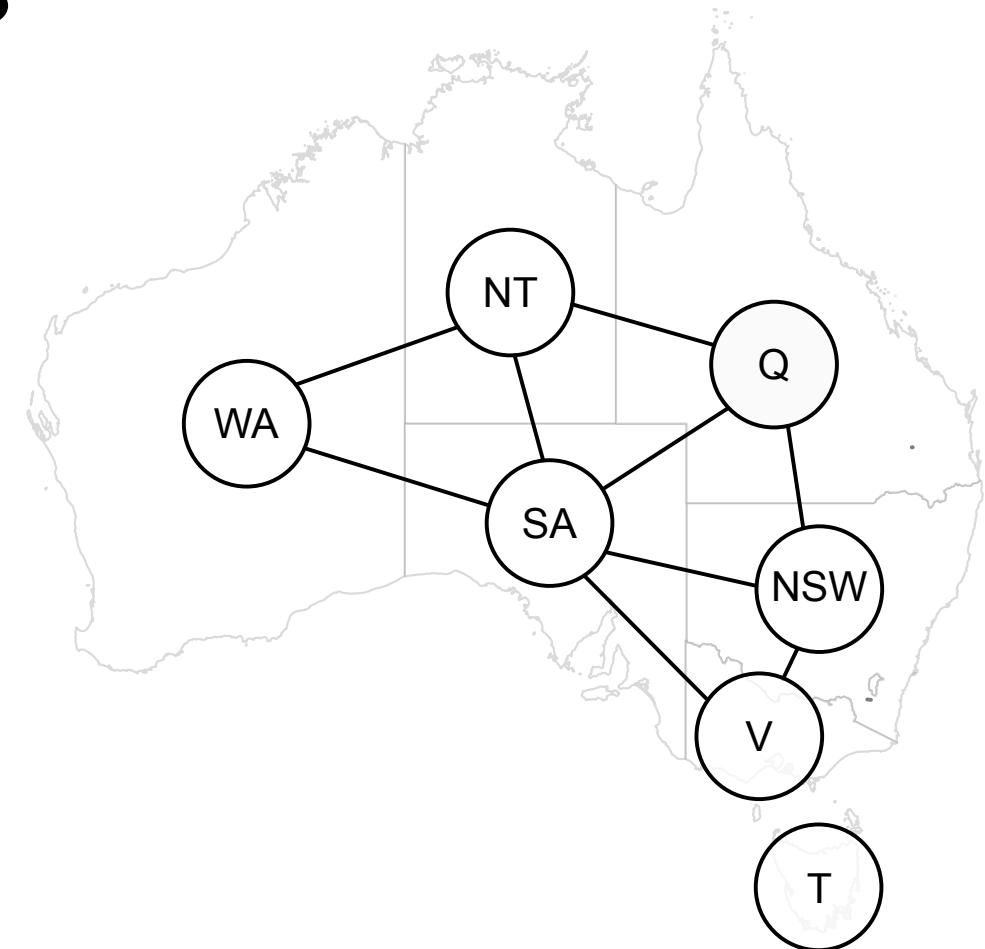
# Example: Map-Colouring

**Variables:** WA, NT, Q, NSW, V, SA, T

**Domains:**  $D_i = \{\text{red, green, blue}\}$

**Constraints:** adjacent regions must have different colours e.g. **WA $\neq$  NT**, etc.

- Constraint graph: nodes are variables, arcs are constraints
- Binary CSP: each constraint relates two variables



# Assignment 2 - CSP

## Part 4: Constraint Satisfaction Problem [5 marks]

This part of the assignment is based on another puzzle called Sudoku. The game consists of 9x9 grid with a few cells filled in with digits. The objective of the game is to fill the remaining cells so that each column, each row, and each of the nine 3x3 sub-grids contain all of the digits from 1 to 9.

The start state of the puzzle is defined as a multidimensional array of numbers, where 0 represents empty cells, for example:

```
grid = [[9, 5, 0, 8, 2, 7, 3, 0, 0],  
        [0, 8, 0, 1, 4, 0, 0, 5, 0],  
        [0, 1, 0, 5, 9, 0, 0, 0, 0],  
        [8, 3, 0, 0, 0, 0, 0, 7, 5],  
        [1, 6, 9, 7, 5, 2, 4, 3, 0],  
        [0, 7, 0, 0, 8, 0, 0, 6, 0],  
        [0, 9, 1, 0, 6, 0, 8, 4, 0],  
        [7, 0, 8, 0, 3, 1, 0, 0, 6],  
        [6, 2, 0, 4, 7, 8, 0, 9, 0]]
```

# Assignment 2 - CSP

```
grid = [[9, 5, 0, 8, 2, 7, 3, 0, 0],  
        [0, 8, 0, 1, 4, 0, 0, 5, 0],  
        [0, 1, 0, 5, 9, 0, 0, 0, 0],  
        [8, 3, 0, 0, 0, 0, 0, 7, 5],  
        [1, 6, 9, 7, 5, 2, 4, 3, 0],  
        [0, 7, 0, 0, 8, 0, 0, 6, 0],  
        [0, 9, 1, 0, 6, 0, 8, 4, 0],  
        [7, 0, 8, 0, 3, 1, 0, 0, 6],  
        [6, 2, 0, 4, 7, 8, 0, 9, 0]]
```

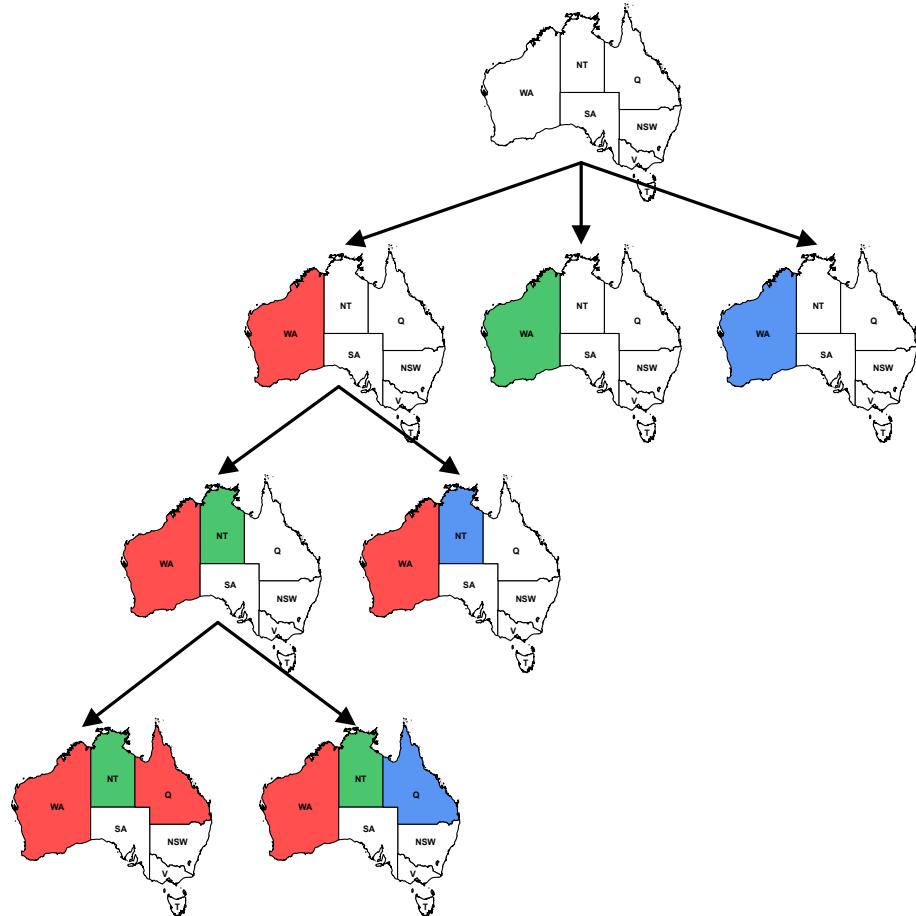
## Task 4.1

Write the code that defines constraint(s) and a function grid\_to\_csp that returns a CSP for a Sudoku puzzle:

```
# define constraint function(s)  
  
# function that returns a csp  
def grid_to_csp(grid):  
    domains = {}  
    constraints = []  
    return CSP(domains, constraints)  
  
# csp  
csp = grid_to_csp(grid)
```

# Constraint Satisfaction Problems

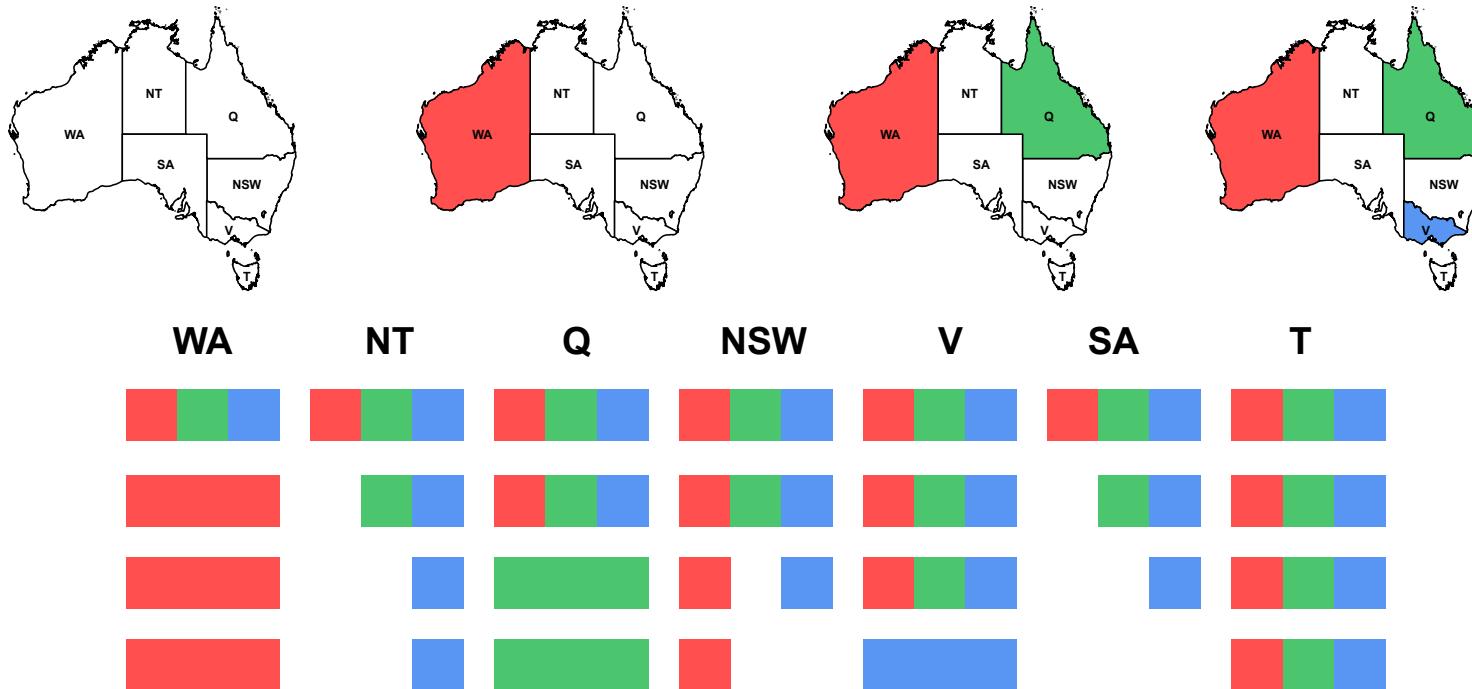
- Backtracking search
- Enhancements to backtracking search
- Local search
  - hill climbing
  - simulated annealing



# Forward checking

Idea:

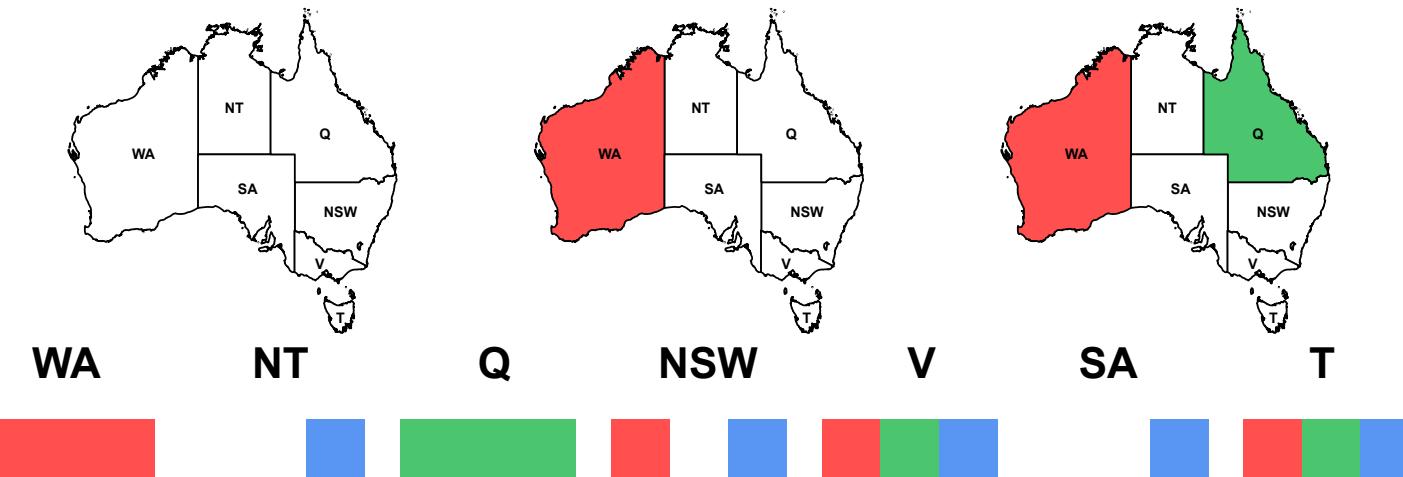
- Keep track of remaining legal values for unassigned variables
- Terminate search when any variable has no legal values
  - prune off that part of the search tree, and backtrack



# Arc consistency

Simplest form of propagation makes each arc consistent

$X \rightarrow Y$  is consistent if *for every* value  $x$  of  $X$  there is *some* allowed  $y$



If  $X$  loses a value, neighbours of  $X$  need to be rechecked.

# Assignment 2 - CSP - Backtracking

## Task 4.2

Write the code that solves Sudoku csp using a search algorithms of your choice, print the solution.

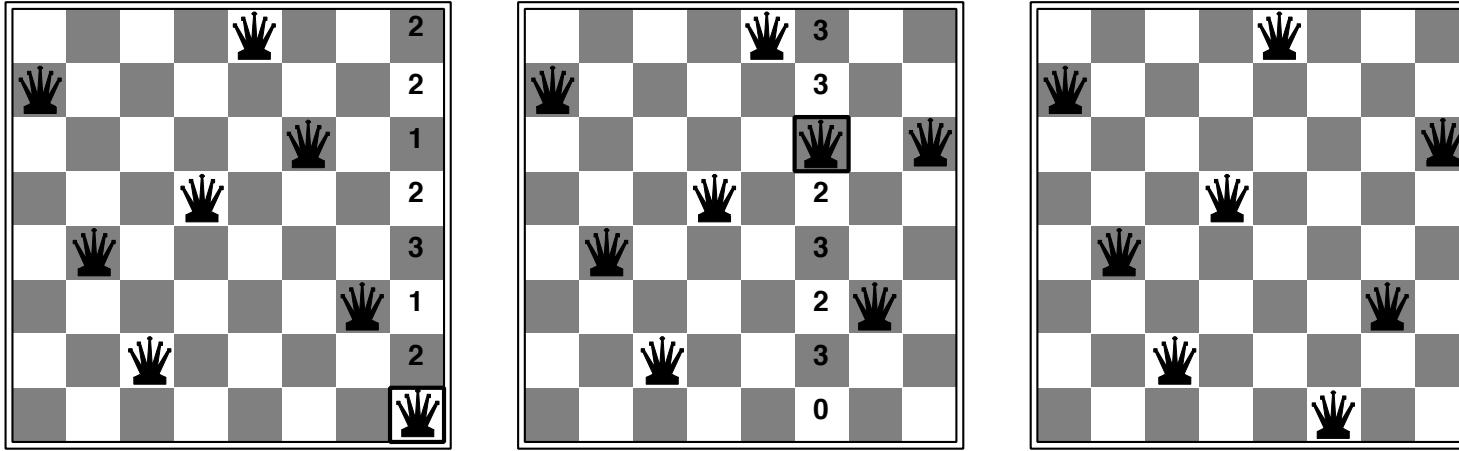
```
# your code
```

## Task 4.3

Describe what do nodes represent in the search tree. Explain your choice of the search algorithm and compare it with any other search algorithm implemented in this assignment.

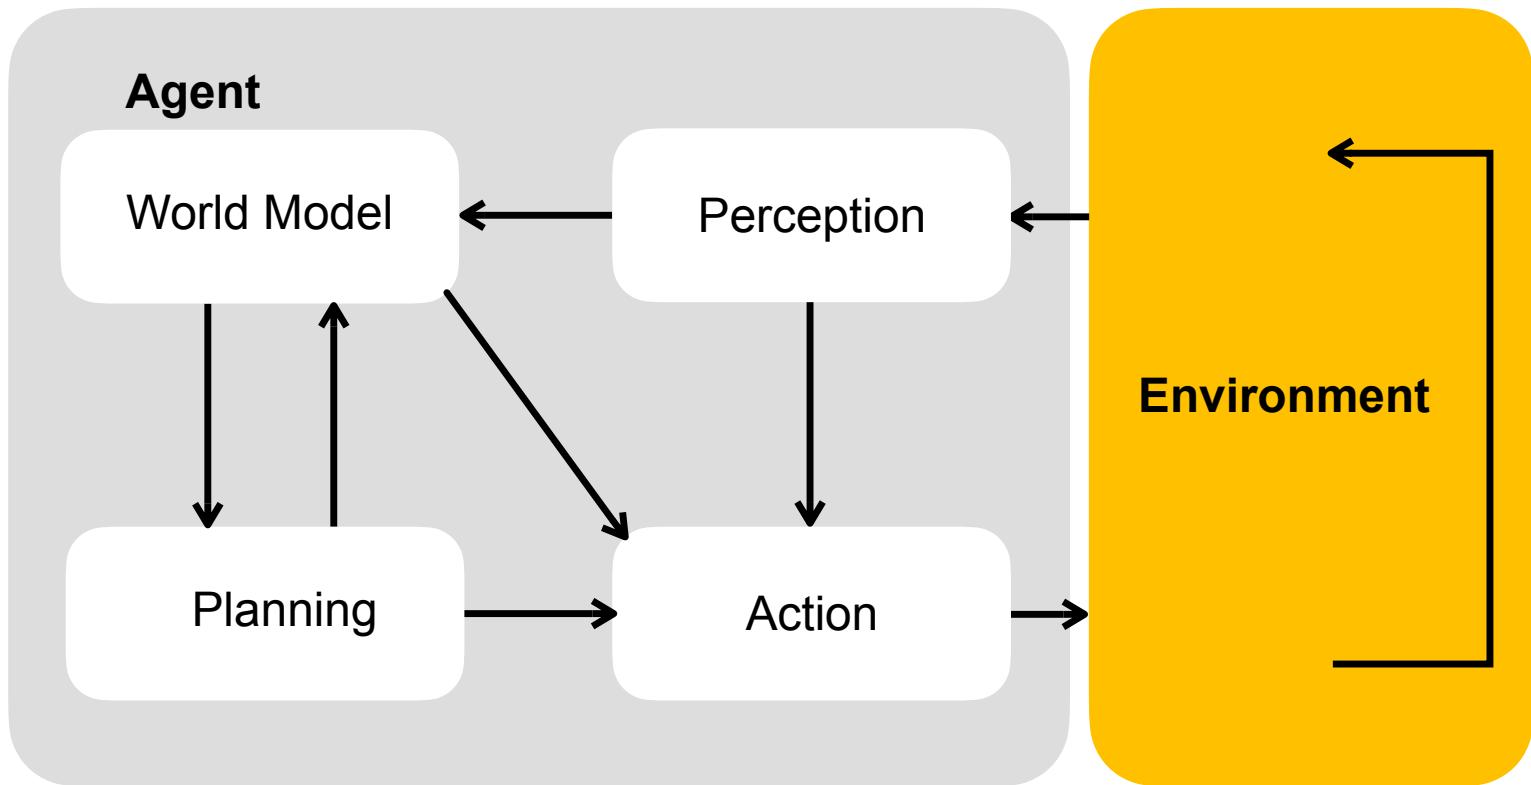
```
# your answer
```

# Hill-climbing by min-conflicts



- Variable selection: randomly select any conflicted variable
- Value selection by **min-conflicts** heuristic
  - choose value that violates the fewest constraints

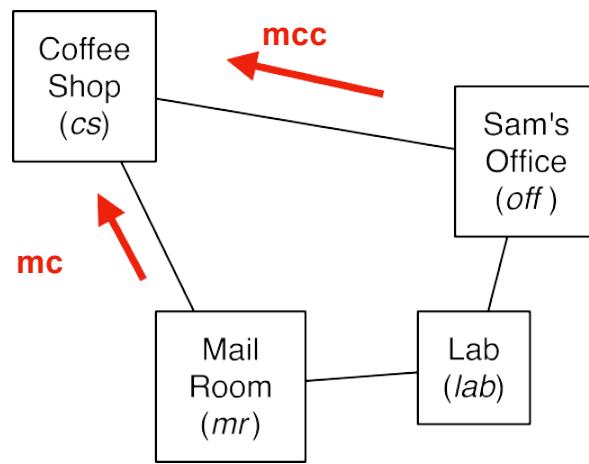
# Planning Agent



# Actions

- A deterministic **action** is a partial function from states to states.
- The **preconditions** of an action specify when the action can be carried out.
- The **effect** of an action specifies the resulting state.

# Delivery Robot Example



Features:

*RLoc – Rob's location*  
*RHC – Rob has coffee*  
*SWC – Sam wants coffee* MW –  
*Mail is waiting*  
*RHM – Rob has mail*

Features to describe states

Actions:

*mc – move clockwise*  
*mcc – move counterclockwise*  
*puc – pickup coffee*  
*dc – deliver coffee*  
*pum – pickup mail*  
*dm – deliver mail*

Robot actions

# STRIPS Representation

- Each action has a:
  - precondition that specifies when the action can be carried out.
  - effect a set of assignments of values to primitive features that are made true by this action.
  - Often split into an ADD list (things that become true after action)
  - and DELETE list (things that become false after action)

Assumption: every primitive feature not mentioned in the effects is unaffected by the action.

# Example STRIPS Representation

Pick-up coffee (puc):

- precondition: [cs,  $\neg$ rhc]
- effect: [rhc]

Deliver coffee (dc):

- precondition: [off, rhc]
- effect: [ $\neg$ rhc,  $\neg$ swc]

# Feature-Based Representation of Actions

- For each action:
  - precondition is a proposition that specifies when the action can be carried out.
- For each feature:
  - causal rules that specify when the feature gets a new value and
  - frame rules that specify when the feature keeps its value.

# Example Feature-Based Representation

- Precondition of pick-up coffee (puc):

$$RLoc=cs \wedge \neg rhc$$

- Rules for “location is cs”:

$$RLoc'=cs \leftarrow Rloc = off \wedge Act=mcc$$

$$RLoc'=cs \leftarrow Rloc = mr \wedge Act=mc$$

$$RLoc'=cs \leftarrow Rloc = cs \wedge Act \neq mcc \wedge Act \neq mc$$

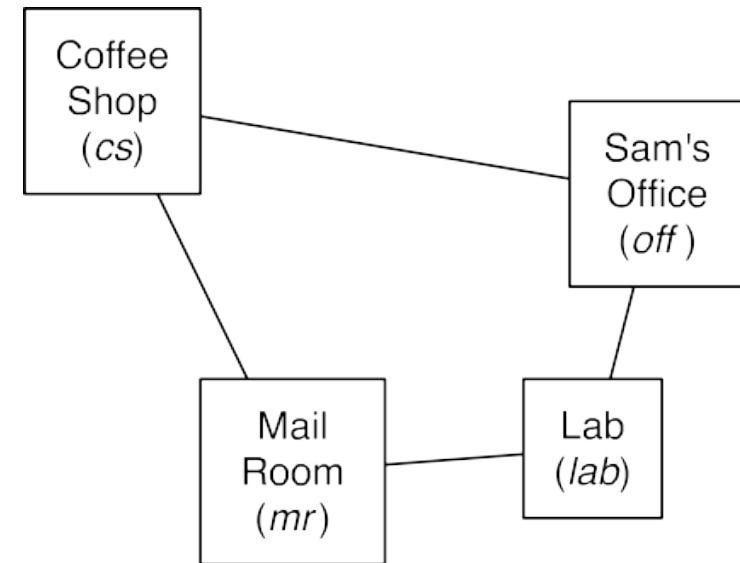
- Rules for “robot has coffee”

$$rhc' \leftarrow Act=puc \wedge \neg rhc$$

$$rhc' \leftarrow rhc \wedge Act \neq dc$$

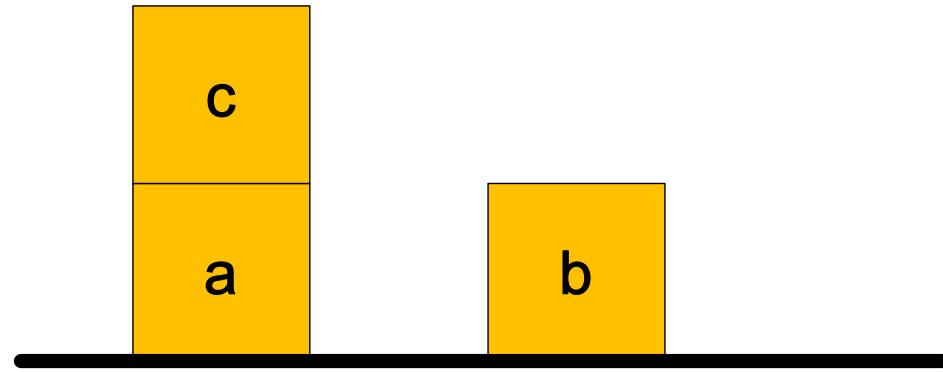
causal rules

frame rule



# Defining Goals and Possible Actions

- Example of goals:  
`on(a,b), on(b,c)`
- Example of action:  
`move(a, 1, b)`  
(Move block a from 1 to b)
- Action preconditions:  
`clear(a), on(a,1), clear(b)`
- Action effects:  
`on(a,b), clear(1), ~on(a,1), ~clear(b)`



“add” (true after action)

“delete” (no longer true after action)  
( $\sim = \neg$ )

# Planning

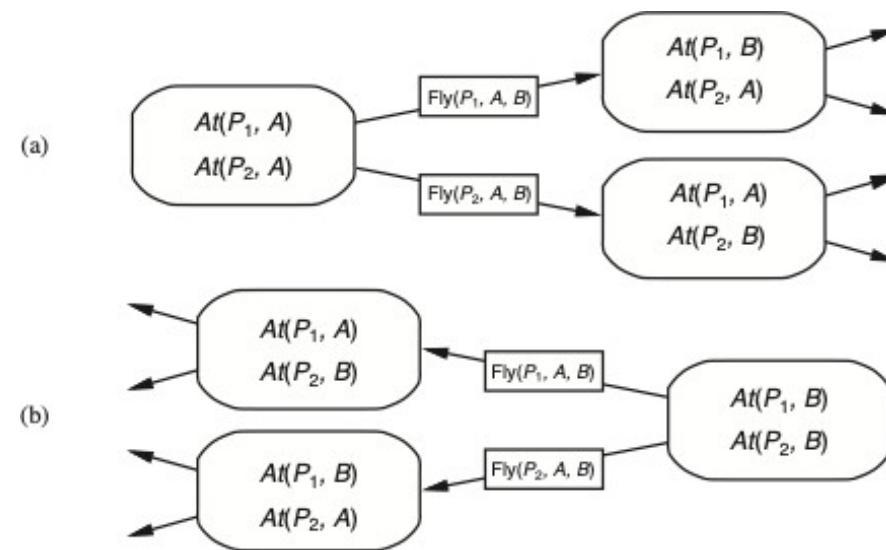
- Plan — sequence (or ordered set) of actions to achieve some goal
- Planner — problem solver that produces plans
- Goal — typically a conjunction of literals
- Initial State — typically a conjunction of literals

Blocks World Example for goal  $on(B, C) \wedge on(C, Table)$

- $move(C, A, Table), move(B, Table, C)$

# Simple Planning Algorithms

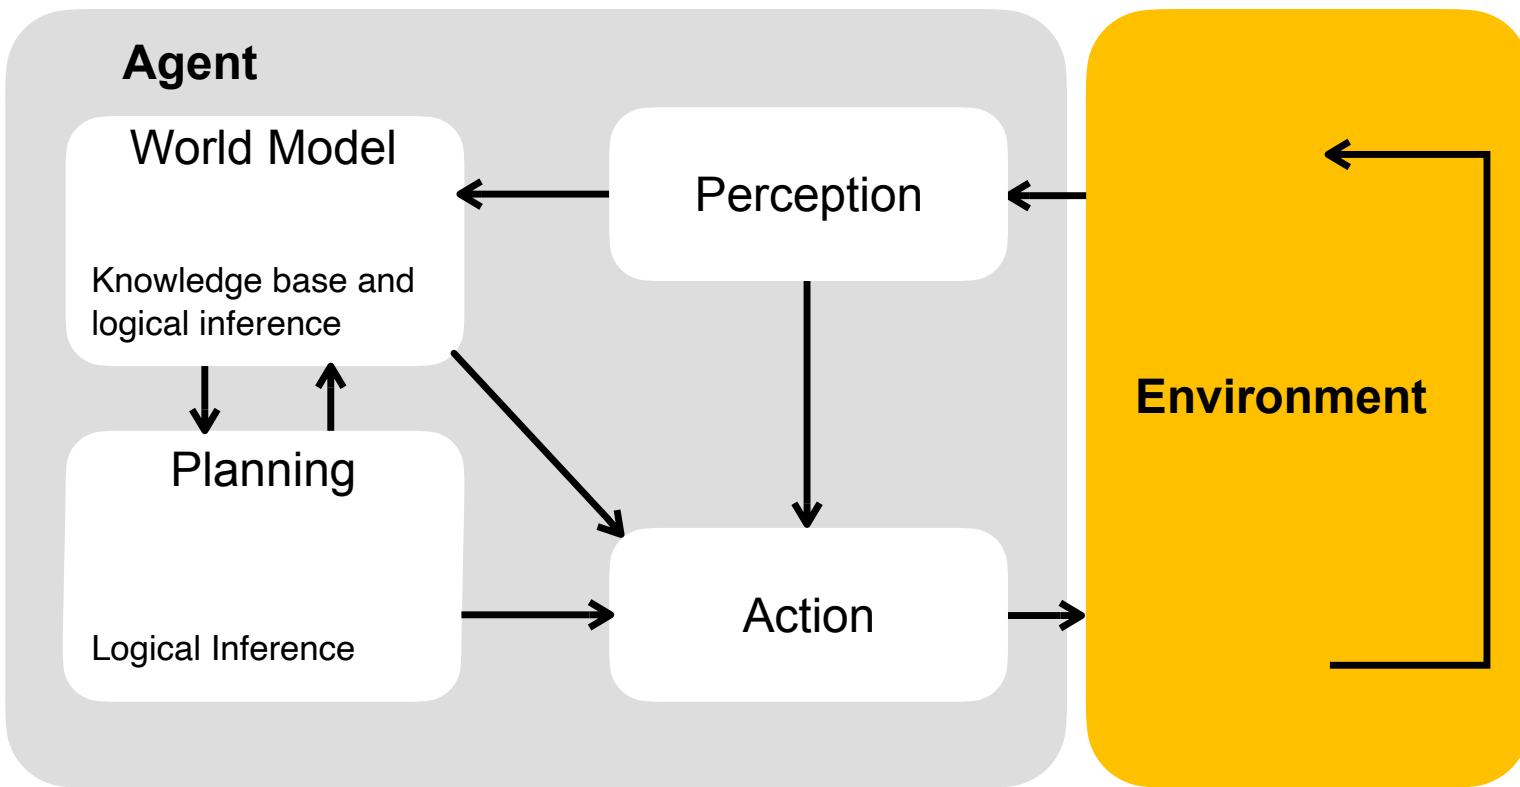
Forward search and goal regression



# Planning

- Planning is the process of choosing a sequence of actions to achieve a goal.
- An action is a partial function from a state to a state.
- Two representations for actions that exploit structure in states are
  - the STRIPS representation, which is an action-centric representation,
  - the feature-based representation of actions, which is a feature-centric representation.
  - The feature-based representation is more powerful than the STRIPS representation; it can represent anything representable in STRIPS, but can also represent conditional effects.

# Models and Planning



# Knowledge Based Agent

The agent must be able to:

- represent states, actions, etc.
- incorporate new percepts
- update internal representations of the world
- deduce hidden properties of the world
- determine appropriate actions

# Propositional Logic: Syntax

Propositional logic is the simplest logic

- “It is not the case that the sky is blue”:  $\neg B$
- (alternatively “the sky is not blue”)
- “The sky is blue and the grass is green”:  $B \wedge G$
- “Either the sky is blue or the grass is green”:  $B \vee G$
- “If the sky is blue, then the grass is not green”:  $B \rightarrow \neg G$
- “The sky is blue if and only if the grass is green”:  $B \leftrightarrow G$
- “If the sky is blue, then if the grass is not green, the plants will not grow”:  $B \rightarrow (\neg G \rightarrow \neg P)$

# Example – Complex Sentence

R	S	$\neg R$	$R \wedge S$	$\neg R \vee S$	$(R \wedge S) \rightarrow (\neg R \vee S)$
True	True	False	True	True	True
True	False	False	False	False	True
False	True	True	False	True	True
False	False	True	False	True	True

Thus  $(R \wedge S) \rightarrow (\neg R \vee S)$  is a **tautology**

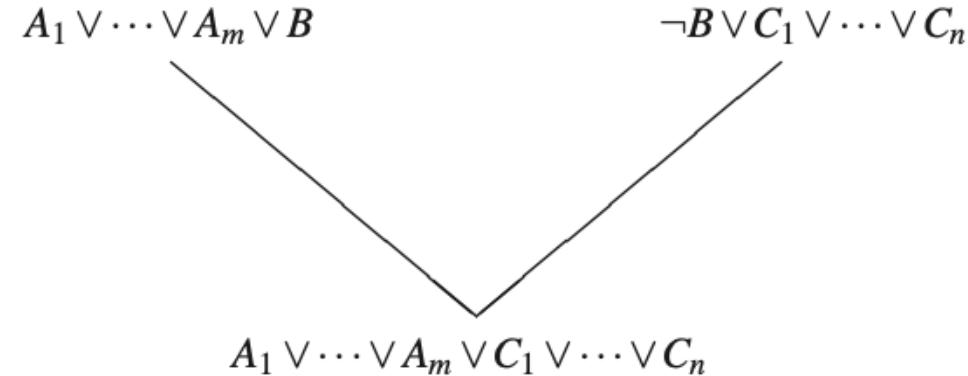
# Propositional Logic

- Use letters to stand for “basic” propositions; combine them into more complex sentences using operators for **not**, **and**, **or**, **implies**, **iff**  $(A \vee B \vee \neg C) \wedge (\neg B \vee D)$
- Propositional **connectives**:

$\neg$	<i>negation</i>	$\neg P$	<i>“not P”</i>
$\wedge$	<i>conjunction</i>	$P \wedge Q$	<i>“P and Q”</i>
$\vee$	<i>disjunction</i>	$P \vee Q$	<i>“P or Q”</i>
$\rightarrow$	<i>implication</i>	$P \rightarrow Q$	<i>“If P then Q”</i>
$\leftrightarrow$	<i>bi-implication</i>	$P \leftrightarrow Q$	<i>“P if and only if Q”</i>

$P$	$Q$	$\neg P$	$P \wedge Q$	$P \vee Q$	$P \rightarrow Q$	$P \leftrightarrow Q$
<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>True</i>	<i>True</i>
<i>True</i>	<i>False</i>	<i>False</i>	<i>False</i>	<i>True</i>	<i>False</i>	<i>False</i>
<i>False</i>	<i>True</i>	<i>True</i>	<i>False</i>	<i>True</i>	<i>True</i>	<i>False</i>
<i>False</i>	<i>False</i>	<i>True</i>	<i>False</i>	<i>False</i>	<i>True</i>	<i>True</i>

# Resolution



Consider  $A_1 \vee \cdots \vee A_m \vee B$  and  $\neg B \vee C_1 \vee \cdots \vee C_n$

- ▲ Suppose both are True
- ▲ If  $B$  is True,  $\neg B$  is False and  $C_1 \vee \cdots \vee C_n$  is True
- ▲ If  $B$  is False,  $A_1 \vee \cdots \vee A_m$  is True
- ▲ Hence  $A_1 \vee \cdots \vee A_m \vee C_1 \vee \cdots \vee C_n$  is True

# Conjunctive Normal Form

In order to apply Resolution, we must first convert the KB into **Conjunctive Normal Form (CNF)**.

This means that the KB is a conjunction of clauses, and each clause is a disjunction of (possibly negated) literals.

$$(A \vee B \vee \neg C) \wedge (\neg B \vee D)$$

# Conjunctive Normal Form

1. Eliminate double negations: rewrite  $\neg\neg P$  as  $P$
2. Eliminate  $\rightarrow$  rewriting  $P \rightarrow Q$  as  $\neg P \vee Q$
3. Use De Morgan's laws to push  $\neg$  inwards (repeatedly)  
Rewrite  $\neg(P \wedge Q)$  as  $\neg P \vee \neg Q$
4. Use the distributive laws to get CNF  
Rewrite  $(P \wedge Q) \vee R$  as  $(P \vee R) \wedge (Q \vee R)$  [for CNF]

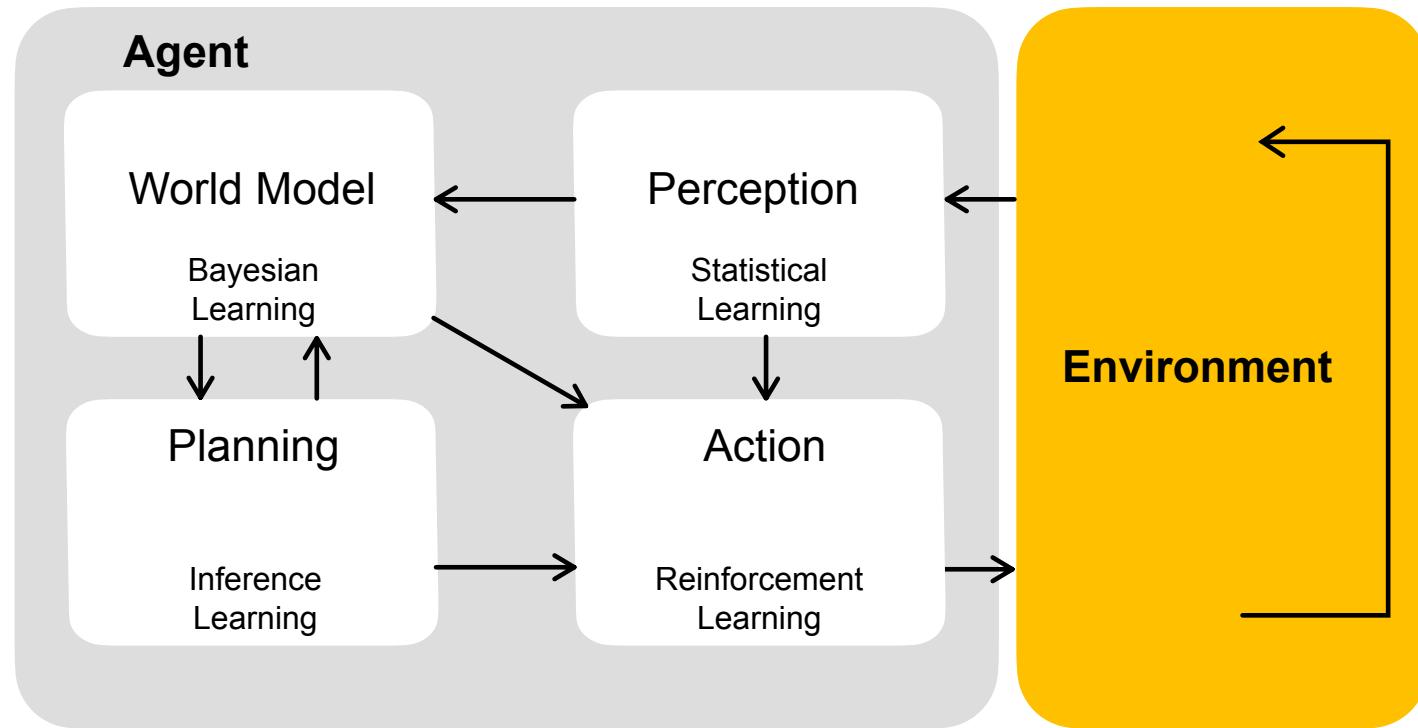
# Resolution by Refutation

- Negate query to be proven (resolution is a refutation system)
- Convert knowledge base and negated query into CNF
- Repeatedly apply resolution to clauses or copies of clauses until either the empty clause (contradiction) is derived or no more clauses can be derived (a copy of a clause is the clause with all variables renamed)
- If the empty clause is derived, answer ‘yes’ (query follows from knowledge base), otherwise answer ‘no’ (query does not follow from knowledge base) . . . and if there are an infinite number of clauses that can be derived, don’t answer at all

# Syntax of First Order Logic

- Objects: people, houses, numbers, theories, colors, football games, wars, centuries ...
- Predicates: red, round, bogus, prime, multistoried, ...  
brother of, bigger than, inside, part of, has color, occurred after, owns, comes between, ...
- Functions: father of, best friend, third inning of, one more than,...

# Learning Agent



# Types of Learning

- Supervised Learning
  - Agent is presented with examples of inputs and their target outputs, and must learn a function from inputs to outputs that agrees with the training examples and generalises to new examples
- Reinforcement Learning
  - Agent is not presented with target outputs for each input, but is periodically given a reward, and must learn to maximize (expected) rewards over time
- Unsupervised Learning
  - Agent is only presented with a series of inputs, and must find and aims to find structure in these inputs

# Supervised Learning

- Given a training set and a test set, each consisting of a set of items for each item in the training set, a set of features and a target output
- Learner must learn **a model** that can **predict** the target output for any given item (characterised by its set of features)
- Learner is given the input features and target output for each item in the training set
  - Items may be presented all at once (batch) or in sequence (online)
  - Items may be presented at random or in time order (stream)
  - Learner **cannot** use the test set **at all** in defining the model
- Model is evaluated by its performance on predicting the output for each item in the **test set**

# Inductive learning

Simplest form: learn a function from examples

$f$  is the target function

An example is a pair  $(x, f(x))$

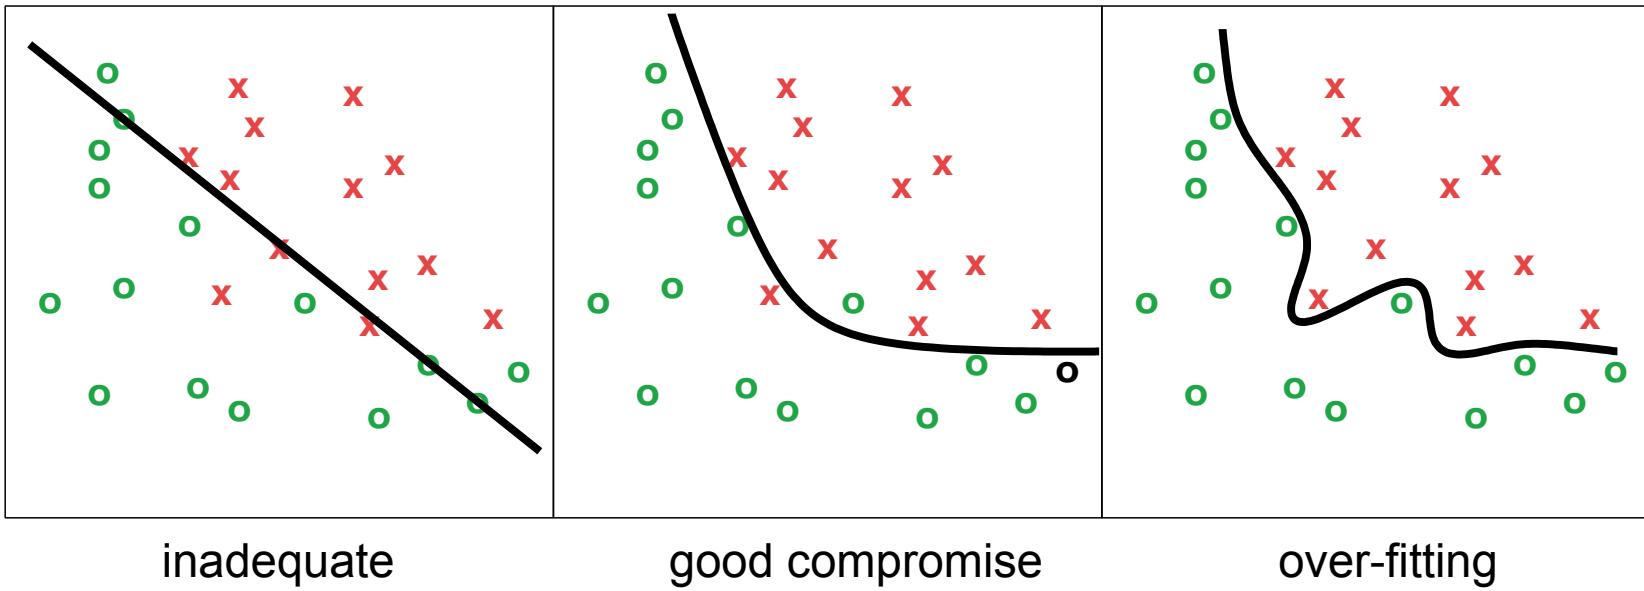
Problem: find a hypothesis  $h$

such that  $h \approx f$

given a training set of examples

# Ockham's razor

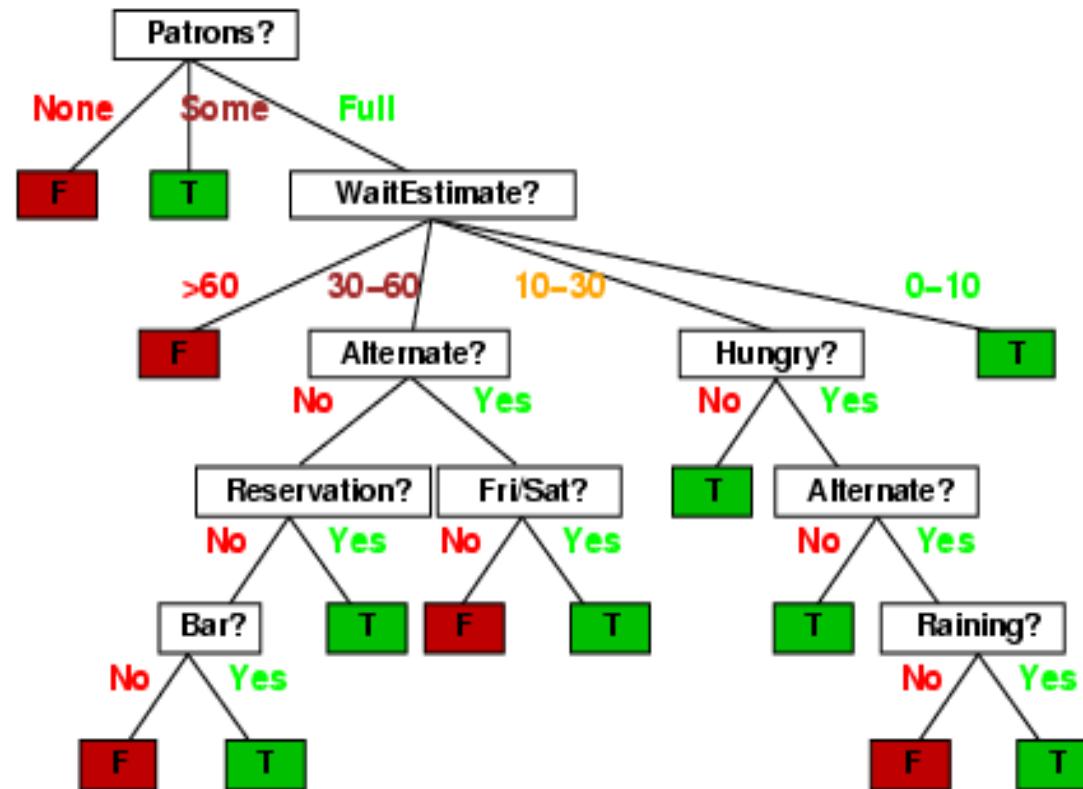
“The most likely hypothesis is the **simplest** one consistent with the data.”



Since there can be **noise** in the measurements, in practice need to make a **tradeoff** between simplicity of the hypothesis and how well it fits the data.

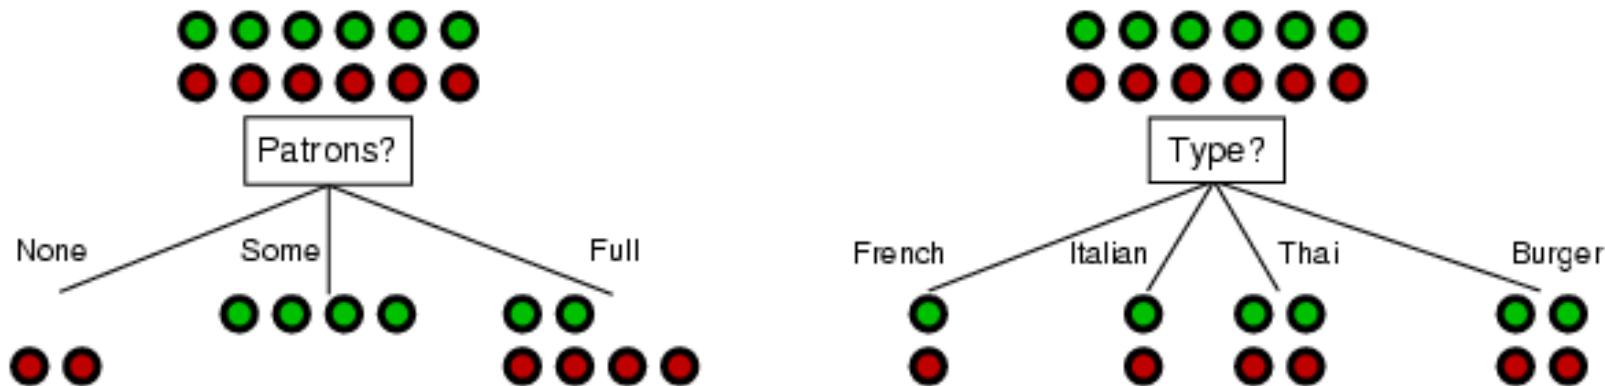
# Decision trees

- One possible representation for hypotheses
  - E.g., here is the “true” tree for deciding whether to wait:



# Choosing an attribute

- Patrons is a “more informative” attribute than Type, because it splits the examples more nearly into sets that are “all positive” or “all negative”.
- An important (or good) attribute splits samples into groups that are (ideally) all positive or negative.



# Entropy

Entropy is a measure of how much information we **gain** when the target attribute is revealed to us. In other words, it is not a measure of how much we know, but of how much we don't know.

If the prior probabilities of the  $n$  target attribute values are  $p_1, \dots, p_n$  then the entropy is

$$H(<p_1, \dots, p_n>) = \sum -p_i \log_2 p_i$$

# Minimal Error Pruning

Should the children of this node be pruned or not?

Left child has class frequencies [7,3]

$$E = 1 - \frac{n+1}{N+k} = 1 - \frac{7+1}{10+2} = 0.333$$

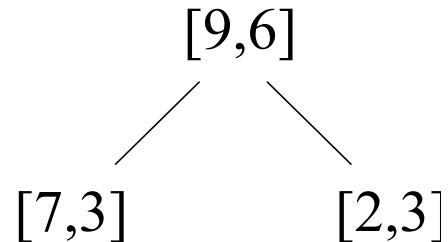
Right child has  $E = 0.429$

Parent node has  $E = 0.412$

Average for Left and Right child is

$$E = \frac{10}{15}(0.333) + \frac{5}{15}(0.429) = 0.365$$

Since  $0.365 < 0.412$ , children should NOT be pruned.



# Probability and Uncertainty

Subjective or Bayesian probability:

Probabilities relate propositions to one's own state of knowledge e.g.  $P(A30|\text{no reported accidents}) = 0.06$

These are not claims of a “probabilistic tendency” in the current situation (but might be learned from past experience of similar situations)

Probabilities of propositions change with new evidence:

e.g.  $P(A30|\text{no reported accidents, 5 a.m.}) = 0.15$

# Prior probability

Prior or unconditional probabilities of propositions

e.g.  $P(\text{Cavity} = \text{true}) = 0.1$  and  $P(\text{Weather} = \text{sunny}) = 0.72$

correspond to belief prior to arrival of any (new) evidence.

Probability distribution gives values for all possible assignments:

$P(\text{Weather}) = \langle 0.72, 0.1, 0.08, 0.1 \rangle$  (normalized, i.e., sums to 1)

# Inference by Enumeration

Start with the joint distribution:

		toothache		~toothache	
		catch	~catch	catch	~catch
cavity	toothache	0.108	0.012	0.072	0.008
	~toothache	0.016	0.064	0.144	0.576

Can also compute conditional probabilities:

$$P(\text{cavity} \mid \text{toothache}) = \frac{P(\text{cavity} \wedge \text{toothache})}{P(\text{toothache})} = \frac{0.016 + 0.064}{0.108 + 0.012 + 0.016 + 0.064}$$

# Conditional Probability

If we consider two random variables  $a$  and  $b$ , with  $P(b) \neq 0$ , then the conditional probability of  $a$  given  $b$  is

$$P(a | b) = \frac{P(a \wedge b)}{P(b)} \quad \text{provided } P(b) > 0 \quad P(a \wedge b) = P(a | b) \cdot P(b) = P(b | a) \cdot P(a)$$

When an agent considers a sequence of random variables at successive time steps, they can be chained together using this formula repeatedly:

$$\begin{aligned} P(X_n, \dots, X_1) &= P(X_n | X_{n-1}, \dots, X_1) P(X_{n-1}, \dots, X_1) \\ &= P(X_n | X_{n-1}, \dots, X_1) P(X_{n-1} | X_{n-2}, \dots, X_1) \\ &= \dots = \prod_{i=1}^n P(X_i | X_{i-1}, \dots, X_1) \end{aligned}$$

# Bayes' Rule

The formula for conditional probability can be manipulated to find a relationship when the two variables are swapped:

$$P(a \wedge b) = P(a | b) \cdot P(b) = P(b | a) \cdot P(a)$$

$$\rightarrow \text{Bayes' Rule } P(a | b) = \frac{P(b | a) P(a)}{P(b)}$$

This is often useful for assessing the probability of an underlying **cause** after an **effect** has been observed:

$$P(\text{Cause} | \text{Effect}) = \frac{P(\text{Effect} | \text{Cause}) P(\text{Cause})}{P(\text{Effect})}$$

# Joint Probability Distribution

We assume there is some underlying joint probability distribution over the three random variables Toothache, Cavity and Catch, which we can write in the form of a table:

		toothache		$\sim$ toothache	
		catch	$\sim$ catch	catch	$\sim$ catch
cavity	0.108	0.012	0.072	0.008	
$\sim$ cavity	0.016	0.064	0.144	0.576	

Note that the sum of the entries in the table is 1.0 .

For any proposition  $\varphi$ , sum the atomic events where it is true:

$$P(\varphi) = \sum_{\{\omega: \omega \models \varphi\}} P(\omega)$$

# Bayesian Networks

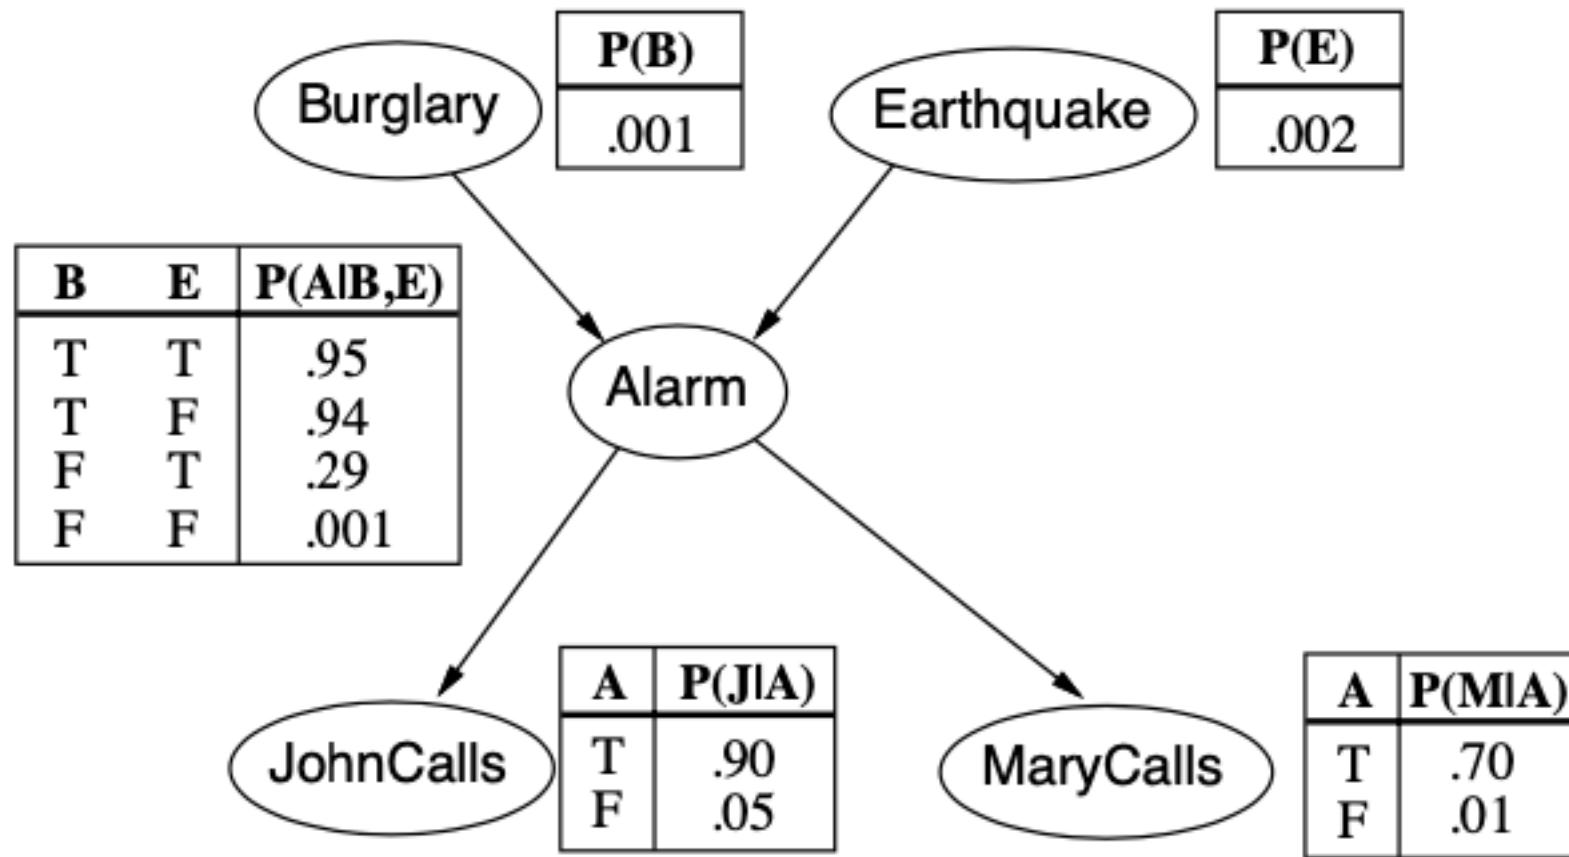
- A **Bayesian network** (also **Bayesian Belief Network**, **probabilistic network**, **causal network**, **knowledge map**) is a directed acyclic graph (DAG) where
  - Each node corresponds to a random variable
  - Directed links connect pairs of nodes – a directed link from node
  - $X$  to node  $Y$  means that  $X$  has a **direct influence** on  $Y$
  - Each node has a conditional probability table quantifying effect of parents on node
- Independence assumption of Bayesian networks
  - Each random variable is (conditionally) independent of its nondescendants given its parents

# Belief Networks

- When there are multiple minimal sets of predecessors satisfying this condition, any minimal set may be chosen to be the parents.
  - There can be more than one minimal set only when some of the predecessors are deterministic functions of others.
- Putting the chain rule and the definition of parents together gives:

$$P(X_1, \dots, X_n) = \prod_{i=1}^n P(X_i | \text{parent}(X_i))$$

# Bayesian Networks



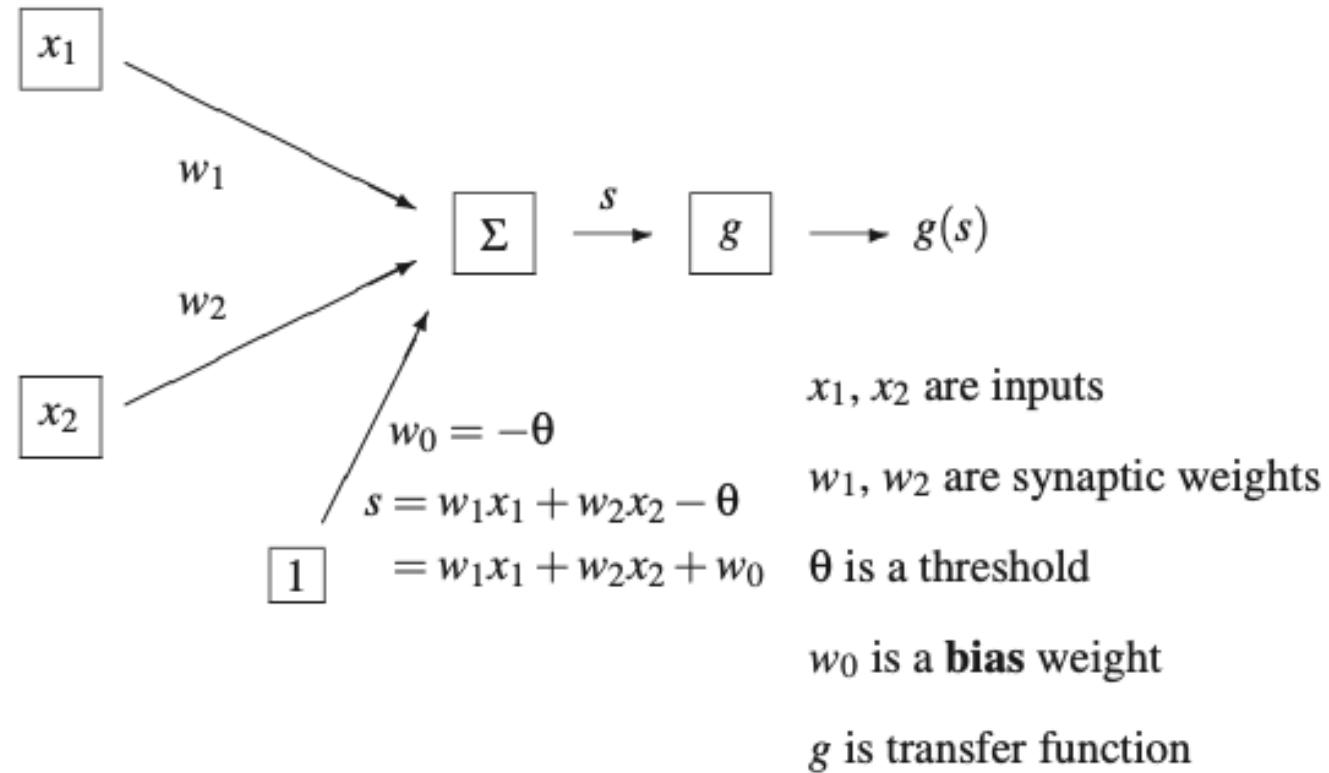
# Artificial Neural Networks

- (Artificial) Neural Networks are made up of nodes which have
  - inputs edges, each with some **weight**
  - outputs edges (with **weights**)
  - an **activation level** (a function of the inputs)
- Weights can be positive or negative and may change over time (learning).
- The **input function** is the weighted sum of the activation levels of inputs.
- The activation level is a non-linear **transfer** function  $g$  of this input:

$$\text{activation}_i = g(s_i) = g\left(\sum_j w_{ij}x_j\right)$$

Some nodes are inputs (sensing), some are outputs (action)

# McCulloch & Pitts Model of a Single Neuron

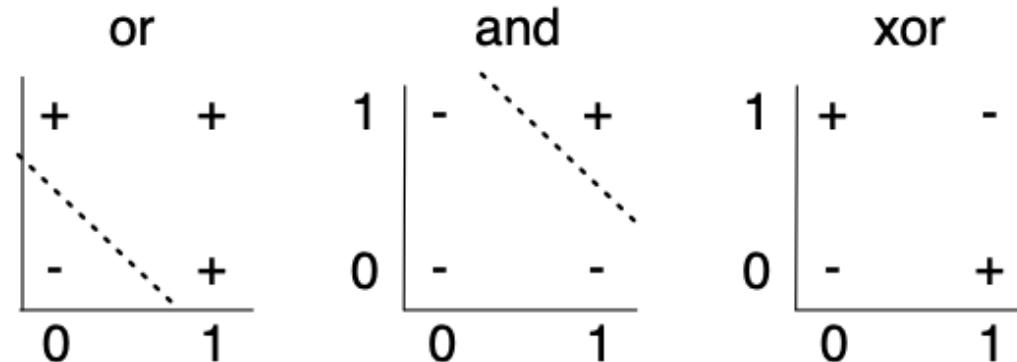


# Linearly Separable

- A classification is **linearly separable** if there is a hyperplane where the classification is true on one side of the hyperplane and false on the other side.
- For the sigmoid function, the hyperplane is when:  $w_0 + w_1X_1 + \cdots + w_nX_n = 0$

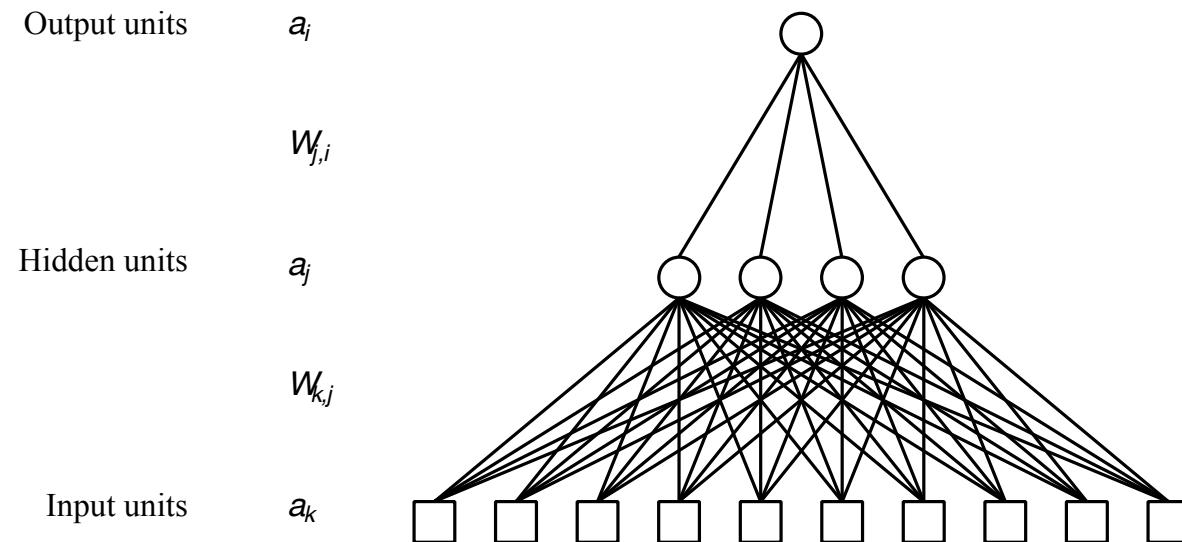
This separates the predictions  $> 0.5$  and  $< 0.5$ .

- Linearly separable implies the error can be arbitrarily small



# Multilayer perceptrons

Layers are usually fully connected;  
numbers of **hidden units** typically chosen by hand



# Gradient Descent

Recall that the error function  $E$  is (half) the sum over all input patterns of the square of the difference between actual output and desired output

$$E = \frac{1}{2} \sum (z - t)^2$$

The aim is to find a set of weights for which  $E$  is very low.

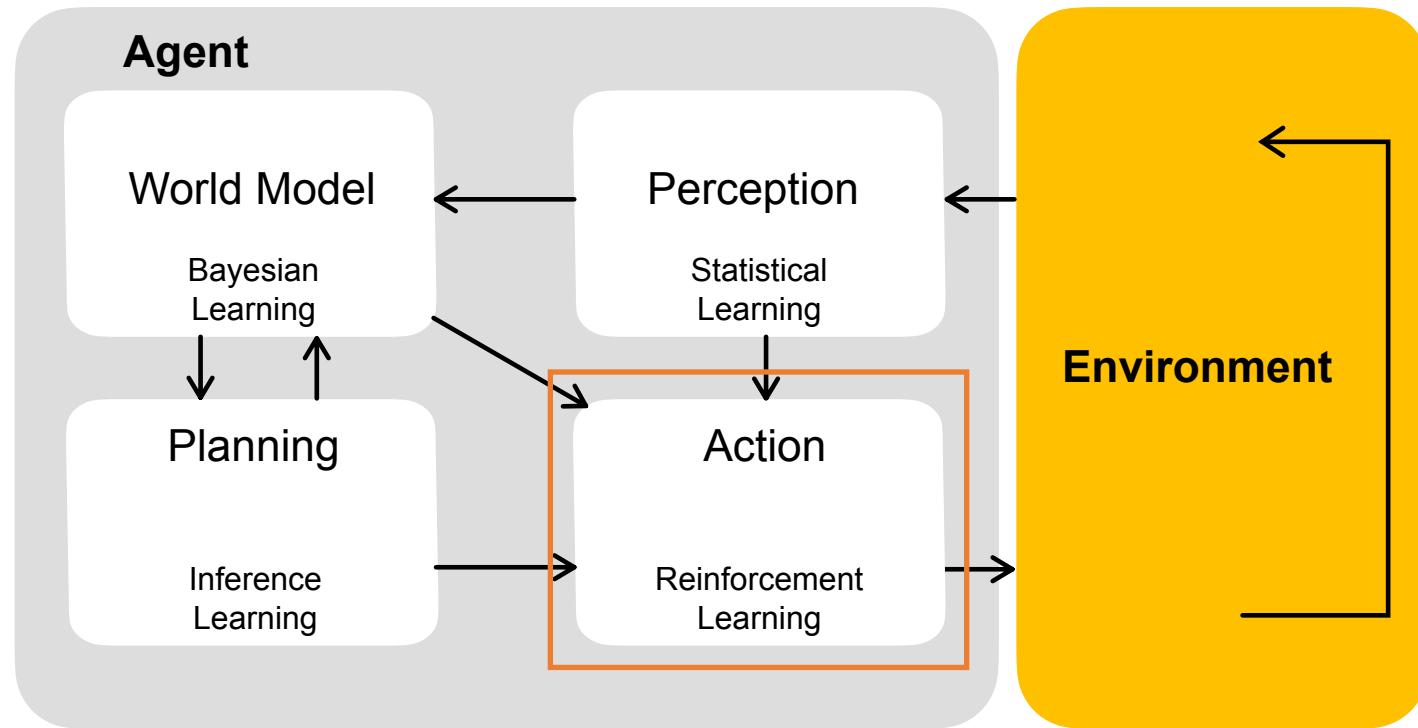
If the functions involved are smooth, we can use multi-variable calculus to adjust the weights in such a way as to take us in the steepest downhill direction.

$$w \leftarrow w - \eta \frac{\partial E}{\partial w}$$

Parameter  $\eta$  is called the learning rate.

- How the cost func effects the particular weight
- Find the weight

# Learning Agent



# States and Actions

- Each node is a *state*
- *Actions* cause transitions from one state to another
- A *policy* is the set of transition rules
  - i.e. which action to apply in a given state
- Agent receives a *reward* after each action
- Actions may be non-deterministic
  - Same action may not always produce same state

# Reinforcement Learning Framework

An agent interacts with its environment.

There is a set  $S$  of *states* and a set  $A$  of *actions*.

At each time step  $t$ , the agent is in some state  $s_t$ . It must choose an action  $a_t$ , whereupon it goes into state

$s_{t+1} = \delta(s_t, a_t)$  and receives reward  $r(s_t, a_t)$ .

In general,  $r()$  and  $\delta()$  can be multi-valued, with a random element

The aim is to find an *optimal policy*  $\pi : S \rightarrow A$  which will maximize the cumulative reward.

# Q-Learning – summary

## $Q$ Value

- How to choose an action in a state?

$$Q(s, a) = r(s, a) + \gamma V^*(s')$$

- The  $Q$  value for an action,  $a$ , in a state,  $s$ , is the immediate reward for the action plus the discounted value of following the optimal policy after that action
- $V^*$  is value obtained by following the optimal policy
- $s' = \delta(s, a)$  is the succeeding state, assuming the optimal policy

# Q-Learning – summary

## *Q* Learning

initialise  $Q(s, a) = 0$  for all  $s$  and  $a$

observe current state  $s$

repeat

    select an action  $a$  and execute it

    observe immediate reward  $r$  and next state  $s'$

$$Q(s, a) \leftarrow r + \max_{a'} Q(s', a')$$

$$s \leftarrow s'$$

# Examination Instructions

- Date and time: Saturday, 06 February 2021 at 2pm (Sydney time)
- Online mode: Moodle
- Reading time: 10 mins
- Time allowed: 2 hours
- This examination counts for 50% of the final mark
- Total number of questions: 36
- Answer all questions

# Examination Instructions

- Questions: multiple choice & open-ended questions
- Questions can have different weight (open-ended questions weight more)
- Similar to the questions from tutorials & assignments
- Some inspired by your homework
- No mark taken for wrong answers
- You can use your personal notes
- **You cannot use any other material (including internet)**
- **Plagiarism policy**
- **Distributing, copying, printing is not permitted.**

# Sample question

Completeness of a search algorithm answers the question:

- (a) Is the algorithm guaranteed to find a solution when there is one?
- (b) Does the strategy find the solution that has the lowest path cost of all solutions?
- (c) How long does it take to find a solution?
- (d) How much memory is needed to perform the search?

# Sample question

Consider this joint probability distribution:

		short		$\neg$ short	
		wide	$\neg$ wide	wide	$\neg$ wide
striped	wide	0.07	0.05	0.08	0.12
	$\neg$ wide	0.14	0.17	0.12	0.25

Compute (to two decimal places):  $\text{Prob}(\text{ short} \vee \neg \text{ wide} \mid \text{ striped})$

- (a) 0.50
- (b) 0.63
- (c) 0.75
- (d) 0.80

# Sample question

Explain why depth-first search is used in the search tree for Constraint Satisfaction Problems. Compare it with breath-first search.

# Beyond COMP3411/9814

- COMP9444 Neural Networks and Deep Learning
- COMP9417 Machine Learning and Data Mining
- COMP4418 Knowledge Representation and Reasoning
- COMP3431 Robotic Software Architecture
- COMP9517 Machine Vision
- 4th Year Thesis topics

# UNSW myExperience Survey

Please remember to fill in the UNSW myExperience Survey.

# Questions?

Good luck!

